

Quash Report

EECS 678

Cole Jurden, Evan Nichols

Overview

Quash (Quite A Shell) uses UNIX system calls to simulate a terminal, its behavior similar to popular shells such as `csh` and `bash`. The program allows all functionality included in a regular shell (background execution, I/O redirection, piping, etc.), along with a selection of built-in functions: `set`, `echo`, `cd`, `pwd` and `jobs`. An explanation of Quash's main execution function is below.

Main Execution Function

The main execution follows a series of `if/else` statements. In `main`, we first check to see if there were other arguments passed in the running of `quash`. If a file is passed in with commands to read from, `quash` reads through that file, splits its contents by `\n`, and executes each line accordingly. If there are no other arguments, `quash` waits until a command is received from the command line and the string containing the command is passed to `parse_command()`. This method first searches the strings for `|`, `&`, `>` and `<`. The results of these `strchr()` calls are stored in variables to be used in the conditionals.

The first conditional relates to the presence of a pipe. If there is a pipe, `parse_command()` splits the command string into two parts: a string of commands before the pipe and after the pipe. Then the program follows standard IPC format, declaring the pipe, `forking` the first process and redirecting `STDOUT`, `execvp`ing the first argument. In the second process, `STDIN` is redirected and the second argument is executed.

If there is no pipe detected, the program moves into an `else` conditional where it checks for other commands. First, it divides the command string using `strtok` with a whitespace delimiter and storing each individual part of the command in a `char*` array. This is mainly because we check the contents of `cmds[0]` for individual commands in our conditionals.

The next conditional check returns true if a command is directed to run in the background. In this case, we remove the `&` from the command, `fork` a child process, and call `parse_command()` using the truncated command string. We add non-completed jobs to our `jobs` array, which is displayed with the `jobs` command.

I/O commands relate to the next two statements. In both cases, we isolate the commands before the `<` or `>` and run them in a child process using the files given after the I/O indicators as redirects for `STDIN` and `STDOUT` respectively. When redirecting, we utilize the `dup2` command. This is also true for pipes.

Custom Commands

Once all these checks have been completed, it is time to look for our redefined commands, `set`, `cd`, `pwd`, `echo` and `jobs`. If the command is `set`, the command arguments is parsed using `strtok` to isolate only the desired path, which is then passed into the `setenv` UNIX system call . If the command is `cd`, the second command argument is passed into the `chdir` UNIX system call, updating the current working directory. If the command is `pwd`, the `getcwd` UNIX system call is invoked and printed to the screen.

Child processes inherit their parents' environmental variables. In order to correctly execute the `echo` function, it was necessary to use two global variables, `HOME_ENV` and `SET_ENV`. Each time the `set` function is called on `HOME` or `PATH`, the `HOME_ENV` or `PATH_ENV` variable are updated to reflect the change. This way, when the user executes `echo $HOME`, the `HOME_ENV` variable is used instead of the `getenv` UNIX system call, which would return the parents' environmental variable .

If the command is `jobs`, we call the function `print_jobs()` . This function utilizes the globally declared job count integer `jc` and the globally declared array of `struct job_t`'s to print current jobs in the same format as `bash` would.

Issues and Bugs

Setenv in a child process

Setting environment variables in a child process will set them for that process, but when an `echo` is called, it will echo the parents' environment variable. While the `setenv` function was being executed correctly, we thought it was wrong as the `echo` command (which initially used the `getenv` UNIX system call), would always return the parents' environmental variable.

Parsing and Cleaning Command Line Input

we initially struggled to parse and store the command line argument in such a way that it could be passed into system calls and `exec` calls without resulting in error. Using the `strtok` C method was an excellent solution, which allows for a string to be broken into tokens (sequences of contiguous characters) separated by delimiters. In our case, the delimiters were " , < , > , = , | ", depending on the command to be executed.

Removal of Jobs

We were not able to effectively remove a job from our `jobs` array. Code for the function `check_jobs()` demonstrates the technique we attempted to implement in principle. Essentially, for all jobs where `waitpid()` returned a value greater than 0, we would decrement job count and iterate through the job array replacing the pid in question. The `check_jobs()` function would be called at the beginning of `parse_command()` and act in the same way that `bash` reports finished background processes.

File Input

Reading input from a file, for certain commands, is functional, although it runs the function three times. We attempted to take the I/O redirection out of a child process entirely, out of one child

process, and in two child processes (within the initial conditional check and then with a recursive call to `parse_command()`) and the only solution that kept quash running was the latter. In attempts to debug we made a very simple program and attempted to incorporate similar functionality. This process yielded the same results we found in the more complicated quash.